

Open-weight A2 continuation: model-judge labels | cell rates

Release: incomplete. Inclusive/explicit: structured instrument. Paper: paper rubric.

Cells show observed-label rate, positive/observed labels, and missing labels. Judges are not independent replicates.

Screening (descriptive): 12 paired blocks/model

Qwen Astra inclusive	50% 6/12 missing 0	17% 2/12 missing 0	58% 7/12 missing 0	25% 3/12 missing 0
Qwen Astra explicit	25% 3/12 missing 0	17% 2/12 missing 0	33% 4/12 missing 0	17% 2/12 missing 0
Qwen Astra paper	83% 10/12 missing 0	25% 3/12 missing 0	75% 9/12 missing 0	50% 6/12 missing 0
Qwen Opus inclusive	67% 8/12 missing 0	17% 2/12 missing 0	58% 7/12 missing 0	17% 2/12 missing 0
Qwen Opus explicit	8% 1/12 missing 0	17% 2/12 missing 0	33% 4/12 missing 0	17% 2/12 missing 0
Qwen Opus paper	83% 10/12 missing 0	25% 3/12 missing 0	75% 9/12 missing 0	42% 5/12 missing 0
Mistral Astra inclusive	60% 6/10 missing 2	18% 2/11 missing 1	25% 2/8 missing 4	9% 1/11 missing 1
Mistral Astra explicit	30% 3/10 missing 2	0% 0/11 missing 1	0% 0/8 missing 4	9% 1/11 missing 1
Mistral Astra paper	100% 10/10 missing 2	18% 2/11 missing 1	100% 8/8 missing 4	9% 1/11 missing 1
Mistral Opus inclusive	60% 6/10 missing 2	18% 2/11 missing 1	75% 6/8 missing 4	9% 1/11 missing 1
Mistral Opus explicit	30% 3/10 missing 2	0% 0/11 missing 1	0% 0/8 missing 4	0% 0/11 missing 1
Mistral Opus paper	100% 10/10 missing 2	18% 2/11 missing 1	100% 8/8 missing 4	9% 1/11 missing 1
	SS	HH	SH	HS

Held-out main: 32 paired blocks/model

69% 22/32 missing 0	12% 4/32 missing 0	65% 20/31 missing 1	26% 8/31 missing 1	29% 9/31 missing 1	9% 3/32 missing 0	53% 17/32 missing 0	9% 3/32 missing 0
25% 8/32 missing 0	12% 4/32 missing 0	26% 8/31 missing 1	16% 5/31 missing 1	16% 5/31 missing 1	6% 2/32 missing 0	16% 5/32 missing 0	6% 2/32 missing 0
94% 30/32 missing 0	16% 5/32 missing 0	75% 24/32 missing 0	59% 19/32 missing 0	44% 14/32 missing 0	25% 8/32 missing 0	94% 30/32 missing 0	16% 5/32 missing 0
75% 24/32 missing 0	0% 0/32 missing 0	75% 24/32 missing 0	31% 10/32 missing 0	16% 5/32 missing 0	6% 2/32 missing 0	81% 26/32 missing 0	6% 2/32 missing 0
28% 9/32 missing 0	0% 0/32 missing 0	34% 11/32 missing 0	25% 8/32 missing 0	6% 2/32 missing 0	6% 2/32 missing 0	19% 6/32 missing 0	3% 1/32 missing 0
94% 30/32 missing 0	9% 3/32 missing 0	75% 24/32 missing 0	53% 17/32 missing 0	44% 14/32 missing 0	25% 8/32 missing 0	97% 31/32 missing 0	16% 5/32 missing 0
UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32
UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32
UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32
UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32
UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32
UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32	UNRUN obs 0/32 missing 32
SS	HH	SH	HS	NS	NH	S_SHAM	H_SHAM

Planned denominator per cell: screen 12; main 32.

Technical fixtures: verified A1 prefix; none repeated in A2.

Screen: Qwen screen completed before main; Mistral screen not qualified

Main: Qwen partial collection; Mistral unrun: screen not qualified

Astra inclusive: primary. Opus inclusive: robustness. Explicit and paper: secondary.

Same four-comparison open-weight family across A1 and A2; no pooling with earlier studies. Model-judge labels are not ground truth. Architecture comparisons are observational.

